

Advanced Computer Networks

Congestion control and scheduling

Prof. Andrzej Duda
duda@imag.fr

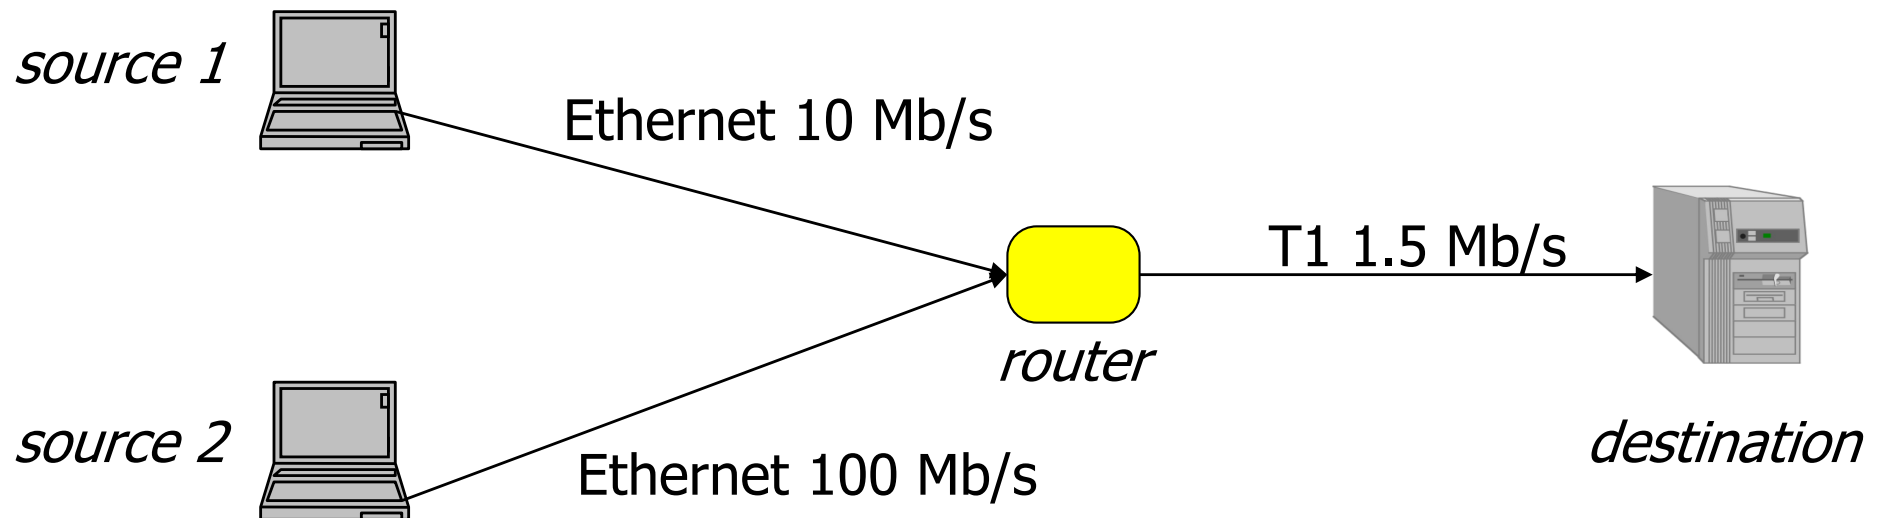
`http://duda.imag.fr`

Contents

- Objectives of Congestion Control
 - efficiency
 - fairness
- Max-min fairness
- Proportional fairness
- Allocation – scheduling disciplines
 - FIFO
 - Priority
 - CBQ
 - Fair queueing
 - DRR



Congestion control

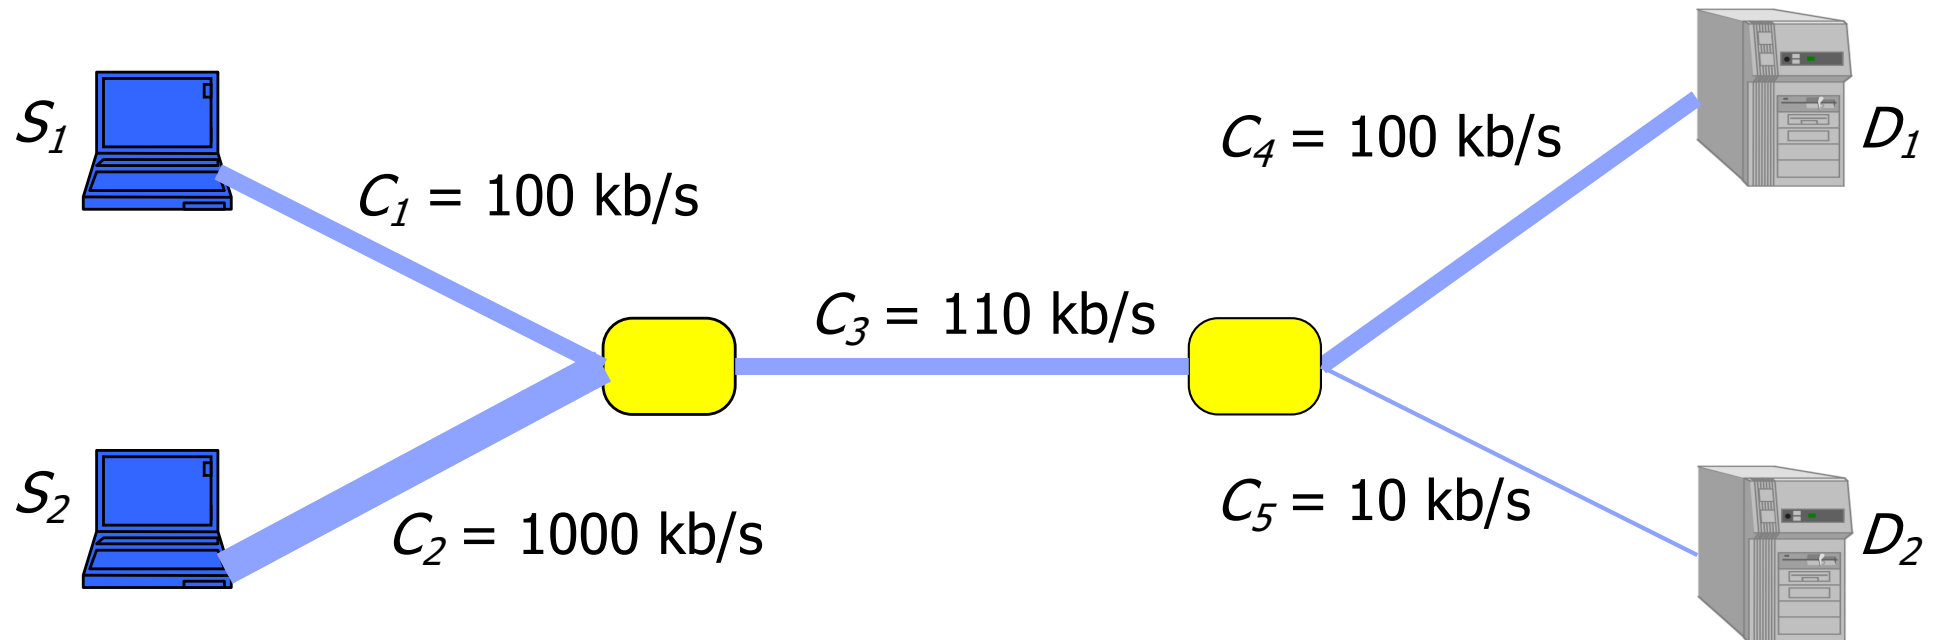


- How to allocate network resources?
 - link capacity
 - buffers at routers or switches
- What to do when the traffic exceeds link capacity?
 - congestion control

Performance criteria

- Efficiency
 - best use of allocated resources
 - max throughput - 100 % utilization
 - min delay - 0 % utilization
- Fairness (équité)
 - fair share to each user
 - different definitions of fairness
 - equal share
 - max-min fairness
 - proportional fairness

Congestion Control - example



- Sources send as much as possible
- Allocation of throughput
 - if the offered traffic exceeds the capacity of a link, all sources see their traffic reduced in proportion of their offered traffic
 - approximately true if FIFO in routers

Throughput allocation

- Throughput x_{ls} : source s on link l
- Traffic λ_s : generated by source s
- Allocation

Our example:

$$x_{11} = \min (\lambda_1, C_1)$$

$$x_{11} = 100$$

$$x_{22} = \min (\lambda_2, C_2)$$

$$x_{22} = 1000$$

$$x_{3i} = \min (x_{ij}, C_3 x_{ij} / (x_{11} + x_{22}))$$

$$x_{31} = 110 \times 100 / 1100 = 10$$

$$x_{32} = 110 \times 1000 / 1100 = 100$$

$$x_{41} = \min (x_{31}, C_4)$$

$$x_{41} = 10$$


$$x_{52} = \min (x_{32}, C_5)$$

$$x_{52} = 10$$

$$\text{throughput } \vartheta = x_{41} + x_{52}$$

$$\text{throughput } \vartheta = 20 \text{ kb/s}$$

Congestion Control - example




- S_1 sends 10 kb/s because it is competing with S_2 on link 3
- S_2 is limited on link 5 anyway

Congestion Control - exemple

- How to increase throughput?
 - if S_2 is aware of the global situation and if it would cooperate
 - S_2 reduces x_{22} to 10 kb/s, because anyway, it cannot send more than 10 kb/s on link 5
 - $x_{31} = 100$ kb/s and $x_{41} = 100$ kb/s without any penalty for S_2
 - throughput is now $\nu = 110$ kb/s

Congestion Control - exemple




- Optimal use of resources

Efficiency criterion

- In a packet network, sources should limit their sending rate by taking into consideration the state of the network. Ignoring this may put the network into congestion collapse
 - network resources are not used efficiently
 - performance indices perceived by sources are not satisfactory
- One objective of congestion control is to avoid such inefficiencies


Efficiency versus Fairness

- Parking lot scenario
 - link capacity: C
 - n_i sources, rate x_i , $i = 1, \dots, I$
 - traffic on link i : $n_0 x_0 + n_i x_i$



Maximal throughput


- For given n_0 and x_0 , maximizing throughput requires that
 - $n_i x_i = C - n_0 x_0$
- Total throughput, measured at the network output
 - $$\begin{aligned} \vartheta &= n_0 x_0 + \sum n_i x_i = n_0 x_0 + \sum (C - n_0 x_0) = \\ &= n_0 x_0 + I(C - n_0 x_0) = IC - (I - 1) n_0 x_0 \end{aligned}$$



$$\vartheta_{\max} = IC \text{ for } x_0 = 0!$$


Maximum throughput

- Example
 - $I = 2, n_0 = n_1 = 1, n_2 = 9$
- The value of x_0 for maximum throughput?
 - 1: C ?
 - 2: $2C$?
 - 3: $0.1 C$?
 - 4: None of the above?



Maximum throughput

- Find x_0 x_1 x_2 such that:
 - $x_0 + x_1 \leq C \rightarrow x_0 + x_1 = C$
 - $x_0 + 9x_2 \leq C$
 - Maximize $x_0 + x_1 + 9x_2 \rightarrow x_0 + x_1 + 9x_2 = 2C$
 - $9x_2 = C$
 - $x_0 = 0, x_1 = C, x_2 = C/9$



Pareto Efficiency (Optimality)


- A feasible **allocation of rates** \mathbf{x}_i is called **Pareto-efficient** iff increasing one source must be done at the expense of decreasing some other source
- For a feasible allocation x_i' , for every i :
- if $x_i' > x_i$ then $x_j' < x_j$
- **Every source has a bottleneck link** (i.e., for every source i there exists a link, used by i , which is saturated)

Pareto Efficiency (Optimality)

- State of resource allocation in which there is no alternative state that would make some people better off without making anyone worse off
- In the case of multiple flows, it means that giving higher rate to a flow cannot reduce the throughput of other flows


Allocation Pareto-Efficient?

- $x_0 = 1, x_1 = 1, x_2 = 2, x_3 = 7$?
- $x_0 = 1, x_1 = 1, x_2 = 4.5, x_3 = 4.5$?
- Both?
- None?
- I don't know?



Pareto-Efficient?

- $x_0 = 0, x_1 = 10, x_2 = 10/9$?
- $x_0 = 0.55, x_1 = 9.45, x_2 = 1.05$?
- $x_0 = 1, x_1 = 9, x_2 = 1$?



Pareto Efficiency


- The Pareto efficient allocations are the ones that **use the resources maximally**
- Maximal efficiency means Pareto optimality.
- Maximizing total throughput is Pareto optimal, but it means shutting down some flows (x_0) this is at the expense of fairness.
- Are there Pareto-efficient allocations that are fair?
What is fairness?
- Egalitarianism (give each flow the same part) is not Pareto-efficient

Fairness

- Maximizing network throughput as a primary objective may lead to large unfairness
 - some sources may get a zero throughput
- Fairness criterion – **equal share to all**
 - let allocate the same share to all sources (egalitarianism), e.g., for $n_i = 1$
 - $x_i = C/2$
 - $\vartheta_{fair} = (I+1)C/2$
 - roughly half of maximal throughput

Fair (equal share)?

- $x_0 = x_1 = x_2 = 0.5$?
- $x_0 = x_1 = x_2 = 1$?
- $x_0 = x_1 = x_2 = 10/9$?




Equal share fairness

- Consider the parking lot scenario for any values of n_i
 - equal share on link i
 - $x_i = C / (n_0 + n_i), i = 1, \dots, I$
 - let decrease x_0 to increase ν^g (we have seen that this maximizes throughput)
 - $x_0 = \min C / (n_0 + n_i),$
 - example
 - $I = 2, n_0 = n_1 = 1, n_2 = 9$
 - link 2: $x_2 = C / (1 + 9) = 0.1 C$
 - link 1: $x_1 = C / (1 + 1) = 0.5 C$
 - $x_0 = \min (0.5 C, 0.1 C) = 0.1 C$
- Allocating equal shares is not a good solution
 - some flows can get more


Example

- Problem
 - link 1: $0.6 C$
 - underutilized
 - link 2: $1 C$



Max-Min Fairness

- We can increase x_1 without penalty for other flows
 - $x_0 = 0.1 C$, $x_1 = 0.9 C$, $x_2 = 0.1 C$
- This allocation is Pareto-efficient!



Max-Min Fairness

- Allocating resources in an equal proportion is not a good solution since some sources can get more than others without decreasing others' shares
- Max-Min fair allocation
 - Min: because of the fairness on bottleneck links
 - Max: because we can increase throughput whenever possible
- For every source i , increasing its rate must force the rate of some other (not richer) source j to decrease
- An allocation is max-min fair if *any rate increase contradicts fairness*
- Max-min fair allocation is Pareto-efficient (converse is not true)

Progressive filling


- Bottleneck link / for source s
 - link / is saturated: $\sum x_i = C$
 - source s on link / has the maximum rate among all sources using that link
- Progressive filling allocation
 - $x_i = 0$
 - increase x_i equally until $\sum x_i = C$
 - rates for the sources that use this link are not increased any more
 - all the sources that do not increase have a bottleneck link (Min)
 - continue increasing the rates for other sources (Max)

Example

- Parking lot scenario
 - $x_i = 0$
 - $x_i = d$ until $n_0 x_0 + n_i x_i = C$
 - bottleneck link for $d_1 = \min (C / (n_0 + n_i))$, source 0 or i
 - $x_0 = \min (C / (n_0 + n_i))$
 - increase other sources
 - $x_i = (C - n_0 x_0) / n_i$
- In our example
 - $x_0 = 0.1 C, x_2 = 0.1 C$
 - $x_1 = 0.9 C$

Max-Min Fair?

- $x_0 = 0$ $x_1 = 10$, $x_2 = 10/9$?
- $x_0 = 1$ $x_1 = 9$ $x_2 = 1$?



Exercise

- $C = 10$
- We have four flows with demands of 2, 2.6, 4, 5
- What is the Max-min allocation to flows?

Exercise

- Two sources 1 and 2 share a capacity link C. The flow x_i of source i is limited by
 - $x_i \leq r_i$ $i = 1, 2$
- Let $C = 9$ Mb/s, $r_1 = 3$ Mb/s, $r_2 = 8$ Mb/s
- Find x_i assuming the allocation is max-min

Proportional Fairness

- Equal share fairness and Max-min fairness
 - per link only
 - do not take into account the number of links used by a flow
 - flows x_o benefit from more network resources than flows x_i
- Another fairness
 - give higher throughput to flows that use less resources
 - give smaller throughput to flows that use more resources
- Proportional fairness

Proportional Fairness

- An allocation of rates x_s is *proportionally fair* if and only if, for any other feasible allocation y_s we have (S sources)

$$\sum_{s=1}^S \frac{y_s - x_s}{x_s} \leq 0$$

- Any change in the allocation must have a negative average change
- Parking lot example with $n_s = 1$
 - max-min fair allocation $x_s = C/2$ for all s
 - let decrease x_0 by δ : $y_0 = C/2 - \delta$, $y_s = C/2 + \delta$, $s = 1, \dots, I$
 - average rate of change is positive - not proportionally fair for $I \geq 2$!

$$\left(\sum_{s=1}^I \frac{2\delta}{c} \right) - \frac{2\delta}{c} = \frac{2(I-1)\delta}{c}$$

Proportional Fairness

- There exists one unique proportionally fair allocation. It is obtained by maximizing

$$J(\vec{x}) = \sum_s \ln(x_s)$$

over the set of feasible allocations for all sources s

Parking lot example

- For any choice of x_0 we should set x_i such that
 - $n_0 x_0 + n_i x_i = C, i = 1, \dots, I$

- Maximize

$$f(x_0) = n_0 \ln(x_0) + \sum_{i=1}^I n_i (\ln(C - n_0 x_0) - \ln(n_i))$$

over the set $0 \leq x_0 \leq C / n_0$.


- The maximum is for

$$x_0 = \frac{C}{\sum_{i=0}^I n_i} \quad x_i = \frac{C - n_0 x_0}{n_i}$$

- If $n_i = 1, x_0 = C / (I+1), x_i = CI / (I+1)$
- Max-min allocation is $C/2$ for all rates - sources of type 0 get a smaller rate, since they use more network resources


Proportionally Fair?

- $x_0 = 1$ $x_1 = 9$ $x_2 = 1$?
- $x_0 = 0.909$ $x_1 = 9.091$ $x_2 = 1.010$?




Comparisons

- $I = 2, n_i = 1$
- max throughput:
 - $x_0 = 0, \text{ throughput} = 2C$
- equal-share and max-min:
 - $x_0 = C/2, x_i = C/2, \text{ throughput} = 1.5C$
- proportional fairness:
 - $x_0 = C/3, x_i = 2C/3, \text{ throughput} = 5C/3$



Scheduling strategies



- Scheduler
 - defines the order of packet transmission
- Allocation strategy
 - throughput
 - which packet to choose for transmission
 - when chosen, packet benefits from a given throughput
 - buffers
 - which packet to drop, when no buffers


FIFO

- Current state of Internet routers
- Allows to share bandwidth
 - proportionally to the offered load
- No isolation
 - elastic flows (rate controlled by the source eg. TCP) may suffer from other flows
 - a greedy UDP flow may obtain an important part of the capacity
 - real time flows may suffer from long delays
- Last packets are dropped - tail drop
 - TCP adapt bandwidth based on losses
- RED (Random Early Detection) techniques
 - choose a packet randomly before congestion and drop it

Priority Queue

- Several queues of different priority
 - source may mark packets with priority
 - eg. ToS field of IP
 - packets of the same priority served FIFO
 - non-preemptive
- Problems
 - starvation - high priority source prevents less priority sources from transmitting
 - TOS field in IP - 3 bits of priority
 - how to avoid everybody sending high priority packets?

Class Based Queueing (CBQ)




- Also called Custom Queueing (CISCO)
- Each queue serviced in round-robin order
- Dequeue a configured byte count from each queue in each cycle
- Each class obtains a configured proportion of link capacity

Characteristics

- Limited number of queues (CISCO - 16)
- Link sharing for Classes of Service (CoS)
 - based on protocols, addresses, ports
- Method for service differentiation
 - assign different proportions of capacity to different classes
 - not so drastic as Priority Queueing
- Avoids starvation

Per Flow Round Robin




- Similar to Processor Sharing or Time Sharing
 - one queue per flow
 - cyclic service, one packet at a time

Characteristics


- It modifies the optimal strategy of sources
 - FIFO: be greedy - send as much as possible
 - RR: use your part the best
 - a greedy source will experience high delays and losses
- Isolation
 - good sources protected from bad ones
- Problems
 - flows sending large packets get more
 - cost of flow classification

Fair Queueing



- Round robin "bit per bit"
 - each packet marked with the transmission instant of the last bit
 - served in the order of the instants
 - allocates rates according to local max-min fairness

Start-Time Fair Queueing



- Round robin "bit per bit"
 - each packet marked with the transmission instant of the first bit
 - served in the order of the instants
 - allocates rates according to local max-min fairness


Deficit Round-Robin (DRR)

- A **quantum** of bits to serve from each connection in order
- Each queue: **deficit counter (dc)** (to store credits) with initial value zero
- Scheduler visits each non-empty queue, compares the packet at the head to dc and tries to serve one quantum of data
 - if $\text{size} \leq (\text{quantum} + \text{dc})$
 - send and save excess in dc: $\text{dc} \leftarrow \text{quantum} + \text{dc} - \text{size}$,
 - otherwise save entire quantum: $\text{dc} += \text{quantum}$
 - if no packet to send, reset dc
- Easier implementation than other fair policies
 - $O(1)$

Deficit Round-Robin

- DRR can handle variable packet size

Quantum size : 1000 bytes



- 1st Round
 - A's dc : 1000
 - B's dc : 200 (served twice)
 - C's dc : 1000
- 2nd Round
 - A's dc : 500 (served)
 - B's dc : 0
 - C's dc : 800 (served)

DRR: performance

- Handles variable length packets fairly
- If weights are assigned to the queues, then the quantum size applied for each queue is multiplied by the assigned weight
- Queues not served during round build up “credits”:
 - only non-empty queues
- Quantum normally set to max expected packet size:
 - ensures one packet per round, per non-empty queue
- Backlogged sources share bandwidth equally
- Simple to implement
 - Similar to round robin

Facts to remember

- In a packet network, sources should limit their sending rate by taking into consideration the state of the network
- Maximizing network throughput as a primary objective may lead to large unfairness
- Objective of congestion control is to provide both efficiency and some form of fairness
- Fairness can be defined in various ways: equal share, max-min, proportional
- Different scheduling algorithms allocate shares of capacity to flows in many ways